

O'REILLY®



Learning AutoML

Automating ML Pipelines with AutoGluon,
Leading Frameworks, and Real-World Integration

Kerem Tomak

Learning AutoML

Automating ML Pipelines with AutoGluon, Leading Frameworks, and Real-World Integration

With Early Release ebooks, you get books in their earliest form—the author’s raw and unedited content as they write—so you can take advantage of these technologies long before the official release of these titles.

Kerem Tomak

O'REILLY®

Learning AutoML

by Kerem Tomak

Copyright © 2025 Kerem Tomak. All rights reserved.

Printed in the United States of America.

Published by O'Reilly Media, Inc., 141 Stony Circle, Suite 195, Santa Rosa, CA 95401.

O'Reilly books may be purchased for educational, business, or sales promotional use. Online editions are also available for most titles (<http://oreilly.com>). For more information, contact our corporate/institutional sales department: 800-998-9938 or corporate@oreilly.com.

Acquisitions Editor: Aaron Black

Development Editor: Shira Evans

Production Editor: Christopher Faucher

Interior Designer: David Futato

Cover Designer: Karen Montgomery

September 2026: First Edition

Revision History for the Early Relesase

- 2025-07-25: First Release

See <http://oreilly.com/catalog/errata.csp?isbn=9798341643185> for release details.

The O'Reilly logo is a registered trademark of O'Reilly Media, Inc. *Learning AutoML*, the cover image, and related trade dress are trademarks of O'Reilly Media, Inc.

The views expressed in this work are those of the author and do not represent the publisher's views. While the publisher and the author have used good faith efforts to ensure that the information and instructions contained in this work are accurate, the publisher and the author disclaim all responsibility for errors or omissions, including without limitation responsibility for damages resulting from the use of or reliance on this work. Use of the information and instructions contained in this work is at your own risk. If any code samples or other technology this work contains or describes is subject to open source licenses or the intellectual property rights of others, it is your responsibility to ensure that your use thereof complies with such licenses and/or rights.

979-8-341-64313-0

Brief Table of Contents (*Not Yet Final*)

Preface (available)

Part 1: Introduction to AutoML

1 What Is Automated Machine Learning (AutoML)? (available)

2 The Rise and Importance of AutoML (available)

Part 2: Understanding Core Concepts and Techniques (unavailable)

3 Hyperparameter Optimization (unavailable)

4 Neural Architecture Search (NAS) (unavailable)

5 Automated Feature Engineering and Data Preprocessing (unavailable)

Part 3: Hands-On with an AutoML Toolkit (unavailable)

6 Working with Tabular Data (unavailable)

7 Processing Text and NLP Data (unavailable)

8 Analyzing Visual Data with Computer Vision (unavailable)

9 Forecasting with Time Series Data (unavailable)

10 Modeling with Multi-Modal Data (unavailable)

Part 4: Integrating AutoML into Production Pipelines (unavailable)

11 Workflow Integration with MLflow and Kubeflow (unavailable)

12 Automating Data Pipelines with Airflow (unavailable)

13 Deployment and CI/CD for AutoML Models (unavailable)

14 Practical Applications Across Industries (unavailable)

15 Comprehensive Case Studies (unavailable)

Preface

A NOTE FOR EARLY RELEASE READERS

With Early Release ebooks, you get books in their earliest form—the author’s raw and unedited content as they write—so you can take advantage of these technologies long before the official release of these titles.

This will be the Preface of the final book. Please note that the GitHub repo will be made active later on.

If you’d like to be actively involved in reviewing and commenting on this draft, please reach out to the editor at sevans@oreilly.com.

Throughout my career spanning decades in data science and analytics, I’ve witnessed a remarkable transformation in how organizations approach machine learning. What once required teams of PhD-level experts and months of painstaking manual work can now be accomplished by domain experts in days or even hours. This democratization of machine learning capabilities through AutoML represents one of the most significant shifts I’ve observed in the field.

Yet despite AutoML’s growing importance, I noticed a critical gap in available resources. Most AutoML documentation focuses on specific tools or provides high-level overviews without practical implementation guidance. Academic papers dive deep into algorithmic details but often lack real-world context. What was missing was a comprehensive resource that bridges theory and practice—one that explains not just how AutoML works, but when and why to use different approaches, and how to integrate them into actual business workflows.

This gap became particularly apparent in my work within organizations across finance, healthcare, retail, and tech. Time and again, I encountered

talented professionals who understood their domain problems deeply but struggled to navigate the AutoML landscape effectively. They needed guidance on selecting the right tools, understanding the trade-offs between different approaches, and implementing solutions that would actually work in production environments - not to mention in regulated industries machine learning models are somewhat still considered black box and frowned upon.

I gained a lot of hands-on experience while searching for a solution as I experimented with all of the tools in this book while building a product for my company med-essence.de. I found AutoGluon to be a fit for our purpose and used it to put one of our use cases in production. While I was able to gather a lot of information across the internet and a variety of papers and books, I had to learn from github pages, fragmented resources and sometimes not so complete documentation of open source software by testing and learning. The idea of this book mostly originated from this multi-year effort.

My recent **collaboration** with Thomas Davenport and Ian Barkin on a topic exploring how AI tools enable “citizen developers” further reinforced this need. We observed that while technology is rapidly democratizing AI development, the gap between available tools and practical knowledge remains substantial. Organizations need more than just access to AutoML platforms—they need a clear understanding of how to use these tools strategically and responsibly. It is also a gray zone how these citizen developed solutions make their way into production.

This book represents my attempt to fill that knowledge gap. Drawing from years of hands-on experience implementing analytics solutions, developing products using traditional and AutoML tools like AutoGluon and TPOT, teaching these concepts to diverse audiences, and observing what works (and what doesn't) in real-world deployments, I've crafted a resource that serves both newcomers seeking to understand AutoML fundamentals and experienced practitioners looking to deepen their implementation skills.

Rather than advocating for any particular tool or approach, this book provides a balanced perspective on the AutoML ecosystem. You'll find

detailed coverage of core concepts like hyperparameter optimization and neural architecture search alongside practical tutorials using AutoGluon, one of the most capable and accessible AutoML frameworks available today. Most importantly, you'll gain the knowledge needed to make informed decisions about when and how to apply AutoML in your specific context.

Who Should Read This Book

This book is designed for anyone who wants to understand and effectively apply automated machine learning, regardless of their current technical background. The content is structured to serve multiple audiences simultaneously, with different readers likely to focus on different sections based on their needs and experience levels.

Data Scientists and ML Engineers will find comprehensive coverage of AutoML's technical foundations, comparative analysis of different optimization approaches, and advanced techniques for customizing automated workflows. Even if you're already experienced with traditional machine learning, this book will help you understand how AutoML can accelerate your work and expand your capabilities. The sections on hyperparameter optimization, neural architecture search, and production integration provide depth that goes beyond typical tool documentation.

Business Analysts and Domain Experts represent perhaps the most important audience for this book. If you understand your organization's data and business problems but lack extensive programming experience, the hands-on sections using AutoGluon will enable you to build sophisticated models with minimal code. The industry-specific examples and case studies will help you identify opportunities to apply AutoML in your domain, while the conceptual explanations ensure you understand what's happening under the hood.

Software Engineers and Application Developers who need to integrate machine learning capabilities into their applications will benefit from the production-focused sections covering deployment, CI/CD, and MLOps

integration. You don't need to become a data science expert to effectively leverage AutoML, but you do need to understand how these systems work and how to integrate them reliably into larger software systems.

Students and Educators in data science, computer science, or related fields will find this book serves as both a comprehensive introduction to AutoML concepts and a practical guide for hands-on learning. The progression from fundamentals through advanced applications, combined with real-world projects and case studies, makes it suitable for both self-directed learning and classroom use.

Business Leaders and Decision Makers should focus on the foundational chapters and industry application sections to understand AutoML's strategic implications. While you may not implement solutions directly, understanding AutoML's capabilities and limitations is crucial for making informed decisions about AI investments and team development.

Consultants and Solution Architects working with multiple organizations will appreciate the broad coverage of different AutoML approaches, the comparative analysis of tools and techniques, and the industry-specific guidance. This book provides the knowledge base needed to recommend appropriate solutions across diverse client contexts.

The book assumes basic familiarity with data analysis concepts but doesn't require deep machine learning expertise. Mathematical concepts are explained intuitively, with technical details provided for those who want them. Code examples use Python and focus on practical implementation rather than complex programming concepts.

Whether you're looking to automate existing machine learning workflows, explore new applications of AI in your organization, or simply understand how AutoML is transforming the data science landscape, this book provides the knowledge and practical guidance you need to succeed.

How This Book is Organized

This book is structured as a comprehensive journey from AutoML fundamentals through advanced implementation and production deployment. The four main parts build upon each other systematically, though experienced readers may choose to focus on specific sections based on their immediate needs.

Part I: Foundations of AutoML establishes the conceptual groundwork you'll need throughout the rest of the book. Chapter 1 introduces AutoML's core concepts, explores why it matters for modern organizations, and examines its role in democratizing AI development. Chapter 2 traces AutoML's evolution through three distinct generations, from early academic tools to today's enterprise-ready platforms. These chapters provide essential context for understanding the current AutoML landscape and making informed tool selections.

Part II: Understanding Core Concepts and Techniques dives deep into the algorithms and methods that power AutoML systems. Chapter 3 provides comprehensive coverage of hyperparameter optimization, the foundational technique underlying virtually all AutoML tools. Chapter 4 explores neural architecture search (NAS), showing how modern systems automatically design optimal neural network architectures. Chapter 5 examines automated feature engineering and data preprocessing, crucial capabilities that often determine model success in real-world applications. Together, these chapters give you the technical foundation needed to understand how AutoML works under the hood and make informed decisions about when and how to customize automated workflows.

Part III: Hands-On with an AutoML Toolkit transitions from theory to practice through extensive tutorials using AutoGluon, one of the most capable and accessible AutoML frameworks available. Each chapter focuses on a specific data type and use case: tabular data (Chapter 6), text and NLP (Chapter 7), computer vision (Chapter 8), time series forecasting (Chapter 9), and multimodal data combining multiple data types (Chapter 10). Every chapter includes both conceptual explanations and complete,

runnable code examples that you can adapt to your own projects. The progression from single-modal to multimodal problems reflects the increasing sophistication of modern AutoML capabilities.

Part IV: Integrating AutoML into Production Pipelines addresses the critical gap between proof-of-concept models and production-ready systems. Chapter 11 covers workflow integration with popular MLOps tools like MLflow and Kubeflow. Chapter 12 focuses on automating data pipelines using Apache Airflow, ensuring your models can operate reliably on fresh data. Chapter 13 provides comprehensive guidance on deployment and CI/CD practices specific to AutoML models. Chapter 14 examines practical applications across major industries, highlighting common challenges and proven solutions. Chapter 15 presents detailed case studies that demonstrate end-to-end AutoML implementations in real-world scenarios.

Reading Strategies for Different Goals:

If you're **new to AutoML**, read the book sequentially from Part I through Part III, then select relevant sections from Part IV based on your deployment needs. This approach ensures you build solid conceptual foundations before diving into implementation details.

If you're an **experienced data scientist** looking to add AutoML to your toolkit, you might skim Part I for context, focus heavily on Part II for technical depth, then jump to specific chapters in Parts III and IV that align with your current projects.

If you're a **business analyst or domain expert** primarily interested in hands-on capabilities, start with Chapter 1 for context, then focus on Part III for practical skills. Return to relevant sections of Parts II and IV as your understanding deepens.

If you're a **software engineer** tasked with integrating AutoML into production systems, concentrate on Chapters 1-2 for context, selected chapters from Part III for understanding AutoML outputs, and Part IV for implementation guidance.

Hands-On Learning Approach:

Throughout the book, especially in Parts III and IV, you'll find a strong emphasis on learning by doing. Code examples are designed to be immediately runnable, with clear explanations of each step. Projects range from beginner-friendly tutorials to complex, multi-component systems that mirror real-world deployments. Each hands-on chapter includes exercises and suggested extensions that encourage you to adapt the techniques to your own data and problems.

The case studies in Chapter 15 deserve special attention, as they demonstrate how the concepts and techniques from earlier chapters come together in complete, end-to-end solutions. These cases are drawn from actual implementations across different industries, providing realistic examples of both successes and challenges you're likely to encounter.

By the time you complete this journey, you'll have both the conceptual understanding and practical skills needed to evaluate, implement, and deploy AutoML solutions effectively in your own context.

Chapter 1. What Is Automated Machine Learning (AutoML)?

A NOTE FOR EARLY RELEASE READERS

With Early Release ebooks, you get books in their earliest form—the author’s raw and unedited content as they write—so you can take advantage of these technologies long before the official release of these titles.

This will be the 1st chapter of the final book. Please note that the GitHub repo will be made active later on.

If you’d like to be actively involved in reviewing and commenting on this draft, please reach out to the editor at sevans@oreilly.com.

When I first entered the field of machine learning, building a model meant spending weeks meticulously selecting and creating features, tuning hyperparameters, and experimenting with different model architectures. This of course happened after a prior round of weeks of cleaning, wrangling and massaging the data in the first place. Each step required specialized knowledge and intuition developed through years of experience.

Throughout my career, I’ve witnessed the field transform dramatically. Machine learning has existed for decades, primarily in the form of supervised learning where models are trained on data with known outcomes to predict future outcomes. What has changed, as Thomas Davenport, one of the most prominent names in data and analytics field, [notes](#), are three critical factors that have created the perfect environment for AutoML’s emergence:

- Growing popularity of ML in organizations

- Use of multiple, sometimes complex algorithms to create a more effective model
- Availability of automated machine learning approaches to make the process more efficient and accessible throughout organizations

In my work with organizations across industries, I've seen firsthand how the explosion of available data, coupled with the need for increasingly granular analyses, has created demand for more sophisticated machine learning applications. At the same time, the supply of qualified data scientists remains limited. Davenport recognizes this fundamental challenge: "The data and algorithms are expanding rapidly, but human capabilities—even those of quantitative professionals and data scientists—are not. Therefore, automating many of the activities in machine learning is perhaps the only way for the supply of analytical capabilities to meet the demand."

Indeed, the evolution of machine learning over the past decade can be characterized by a single word: acceleration. Machine learning projects that once took my teams months to complete can now be finished in days or even hours. This increased speed is not due to an increase in the number of data scientists but rather to the automation of many of the tasks involved in machine learning.

Throughout this book, I'll explore this transformation and address an important question: How much of the machine learning pipeline can be effectively automated? The answer to this question is constantly evolving as AutoML technologies advance, but the implications are significant for businesses, researchers, and practitioners alike.

The Growing Demand for Machine Learning Solutions

The adoption of AI and machine learning (ML) has grown from niche experiments to a core strategic priority across industries. Surveys indicate that **over half of all companies are now using AI in at least one function,**

and more than 77% are exploring AI technologies. Among leading organizations, the engagement is even higher – about 90% of top enterprises have active AI initiatives as of 2023. This broad enthusiasm is driven by tangible gains; companies at the forefront of AI report significantly higher business performance (e.g. 1.5X revenue growth compared to peers), confirming that ML solutions can unlock competitive advantages in efficiency, customer insight, and innovation.

Several market indicators illustrate the surging demand for ML capabilities. Global spending on AI is climbing at over 35% annually, with the overall AI market projected to reach \$730+ billion by 2030. In specific sectors, investment trajectories are striking – for example, AI applications in healthcare are expected to grow from ~\$20B in 2023 to \$188B by 2030 as organizations seek data-driven improvements in diagnostics and patient care. Similar trends are seen in finance (for algorithmic trading, risk modeling, fintech services) and manufacturing (for predictive maintenance, quality control), reflecting a broad-based appetite for ML-driven solutions. This demand is not just among tech companies; traditional industries are increasingly recognizing that leveraging their data with ML is key to staying competitive.

TIP

Projected global AutoML market size, in USD billions (2023–2030). The AutoML segment is expected to grow rapidly (40%+ CAGR) as organizations seek scalable ML solutions without requiring large expert teams.

Crucially, organizations are not just experimenting – they are seeking to operationalize ML at scale. A majority (72%) of businesses have reported integrating AI into routine processes, and most plan to increase their AI investments further. However, building custom ML models from scratch can be slow and costly, prompting interest in more efficient solutions. A study in manufacturing found that using off-the-shelf analytic/ML tools yielded a median ROI of 140%, significantly higher than the 104% ROI from fully custom-built models. In other words, businesses can achieve

faster returns by adopting ready-made ML platforms and automation. This economic incentive, combined with competitive pressure to realize AI benefits quickly, fuels the rise of AutoML as a key approach to meeting the demand. AutoML promises to accelerate ML solution delivery and broaden access, which is increasingly attractive given the growing volume of AI use cases companies want to tackle.

Real-world outcomes underscore why demand keeps rising. For instance, an EY case study described a manufacturer that deployed data analytics (including ML models) on the factory floor and achieved nearly a 40% increase in product yield along with a 60% reduction in unplanned downtime. In finance, **banks using ML for credit scoring and fraud detection have seen measurable improvements – one global bank reported that ML-driven risk models (developed faster with automation) allowed them to increase loan volumes without raising losses, adding \$7–8 million AUD in annual profit.** These successes drive home the point that effective ML solutions can translate directly into ROI. Consequently, demand is rising not only for more ML models, but for faster, more cost-effective ways to develop them. This is the backdrop against which AutoML has emerged as a pivotal technology: organizations need far more ML solutions than their current teams can produce, and they are eager for tools that can automate and accelerate the ML development cycle.

Addressing the Data Science Talent Gap

The rush to adopt machine learning has created a severe talent gap – there are simply not enough trained data scientists to meet industry needs. Nearly every major company is hiring for data science roles, but supply falls short. **A Deloitte analysis projected a shortfall of about 250,000 data scientists in the U.S. by 2024. This scarcity is evident in hiring metrics: data science and analytics jobs take around 45 days to fill on average, about 5 days longer than other tech roles.** The difficulty in recruiting skilled ML experts leads to project backlogs and higher personnel costs, impeding organizations' data

initiatives. In short, the demand for ML solutions is growing (as seen above), yet the limited availability of experts has become a bottleneck.

Moreover, even when companies do have data science teams, those teams spend a huge portion of time on repetitive, lower-level work. **Studies find – and the author confirms this from his own decades long experience in the field! - that data scientists devote ~70–80% of their time to mundane tasks like data cleaning, feature engineering, and tuning algorithms.** This is sometimes called the “80/20 rule” – only 20% of their effort goes into the high-value creative aspects of modeling, while the rest is housekeeping that could potentially be automated. From an efficiency standpoint, this is a poor utilization of scarce (and expensive) expert talent. It has spurred interest in tools that can automate the routine portions of ML work, freeing up data scientists to focus on strategy and domain-specific problem solving.

AutoML has stepped in as a compelling part of the solution to the talent gap. By automating many stages of the ML pipeline (from data preprocessing and feature selection to model selection and hyperparameter tuning), AutoML platforms allow a single data scientist to deliver far more models in the same time than they could via hand-coding. For example, AutoML systems can rapidly test dozens of algorithms and parameter combinations in parallel – something even a large team could only do through arduous manual effort. Early adopters of data science automation report significant time savings: one survey noted that more than 40% of data science tasks were expected to be automated by 2020, yielding increased productivity and broader use of analytics by non-specialists. With automation handling the grunt work, organizations can leverage their limited experts more effectively, and even relatively small teams can undertake complex ML projects.

Beyond making individual data scientists more productive, AutoML helps address the talent gap by enabling a broader pool of people to contribute to ML solution development. This is sometimes described as empowering “citizen data scientists” (employees who aren’t formally trained in ML, but have domain knowledge and some analytical skills). In a **recent co-authored article**, I argue that recent surge in AI tools also enable citizen developers to

take over a large swath of applications by building them themselves. We will discuss democratization in the next section, but it's directly relevant to the talent shortfall: instead of trying to hire an army of PhD-level experts (which the market cannot supply), companies can use AutoML tools to extend ML capabilities to existing staff in analytics, IT, or subject-matter departments. In effect, AutoML is a force-multiplier for the scarce experts and a bridge for skilled generalists to fill the gap.

From a business perspective, closing the talent gap with AutoML and related strategies is crucial. Organizations that cannot fill their data science openings (or cannot afford large ML teams) risk falling behind. By adopting AutoML, these organizations can still progress in AI initiatives – for example, a regional bank without a large data science division could use an AutoML platform to develop credit risk models with existing data analysts, avoiding the need to hire dozens of new experts. Likewise, a marketing firm can upskill its analysts with AutoML to build customer propensity models internally. This not only helps execute projects in the face of talent scarcity, but also reduces reliance on expensive consulting or outsourcing. In summary, AutoML addresses the talent gap on two fronts: automation of labor-intensive ML tasks (so each data scientist is more productive and less overwhelmed) and democratization of ML development (so the circle of people who can build models expands beyond the few hard-to-hire experts).

Democratizing AI Development

One of the most important impacts of AutoML is its role in democratizing AI – making machine learning development accessible to a much wider range of people. Traditional ML development required deep expertise in algorithms, tuning, and coding, effectively limiting it to PhD data scientists and ML engineers. AutoML changes that paradigm by abstracting complex workflows into easier interfaces. As a result, professionals who are not ML specialists (e.g. business analysts, engineers, doctors, marketers) can now create and deploy models with minimal code or even through visual tools.

In other words, AutoML is turning many non-programmers into productive “citizen data scientists.”

In our article “**We are all Programmers Now**”, I co-authored with Davenport and Barkin, we explore how generative AI tools are transforming the workplace by enabling employees without formal coding backgrounds to become “citizen developers.” These tools allow users to create applications through natural language prompts, effectively democratizing software development. This shift challenges the traditional dynamic between IT professionals and end-users, as non-technical staff can now build solutions independently, potentially accelerating innovation and reducing reliance on overburdened IT departments. We emphasize the need for organizations to support this change by providing appropriate training, governance, and integration strategies to harness the benefits while maintaining control and security. This includes AutoML solutions which pre-seeded the genAI boom.

Modern AutoML platforms emphasize user-friendly design to empower these non-experts. Many offer drag-and-drop model builders, guided wizards, or simple point-and-click interfaces to ingest data, select target outcomes, and then let the system train numerous candidate models automatically. For example, Google’s Cloud AutoML and Microsoft’s Azure AutoML allow users to upload a dataset and specify an objective (like predicting churn or classifying images) through a web UI; the service handles all the model training behind the scenes. Similarly, H2O.ai’s Driverless AI provides an interface where an analyst can drop in data and receive an optimized model along with explanations, without writing complex code. By removing the steep learning curve of ML, AutoML enables domain experts to leverage AI directly. A marketing manager can build a customer segmentation model, or a physician can develop a risk prediction model, all without needing a data science team on hand.

Empirical studies confirm this trend of widening participation. AutoML has been called the driving force of “democratizing data science”. Longitudinal research in healthcare informatics, for example, shows AutoML enabling clinicians to create accurate diagnostic models on their own, matching or

even surpassing expert-built models in certain tasks. By automating algorithm selection and optimization, AutoML systems have achieved performance on par with seasoned data scientists in domains like medical image analysis, disease prediction, and more. The key is that these systems let domain professionals experiment with ML without needing extensive ML knowledge. This broadens AI accessibility in practice – hospitals, small businesses, and government agencies can implement ML solutions driven by staff who understand the problem domain, while relying on AutoML’s automation to handle the technical ML details.

Democratization through AutoML is also helping organizations fill the skills gap internally. Rather than hiring outside experts for every project, companies are training their existing workforce to use AutoML tools. For instance, an insurance company might train its actuaries (who have strong data and business knowledge but limited ML background) to use an AutoML platform for building fraud detection models. Those actuaries can then prototype and deploy models themselves, something that would have been impossible a few years ago without a data science degree. **A recent survey noted that citizen data scientists using AutoML and self-service analytics are on the rise, and Gartner famously predicted that they would produce a greater volume of analytics than professional data scientists, due in large part to AutoML automating many tasks.** While that prediction was bold, it underscores a real shift: AutoML is empowering a new class of users to partake in AI development, thereby democratizing who can create value from ML.

It’s important to note that democratization doesn’t mean data scientists become obsolete; rather, it augments collaboration between experts and non-experts. AutoML allows business users to tackle “long tail” analytics problems on their own, or to rapidly prototype an idea, which the data science team can later review and refine. This frees the experts to focus on the most complex or novel challenges (where their specialized skills are truly needed) while enabling quicker turnaround for the more routine models. In practice, many companies adopt a hybrid approach: auto-generated models are vetted by experienced data scientists, or experts set up

the AutoML environment and guardrails (like specifying certain algorithms or fairness constraints), and then domain teams iterate within those guidelines. The net effect is a much greater scale of AI development – more models built and deployed – because the organization is no longer limited only to what the core data science team can hand-craft. AutoML lowers the entry barrier so that AI development becomes a team sport across the enterprise, not the exclusive domain of a few specialists.

Several case studies highlight this democratization in action. In one example, a retail company’s marketing analyst with no formal ML training used a cloud AutoML service to create a customer churn prediction model. The platform handled the heavy lifting (data preprocessing, model training, validation) and produced a model that identified at-risk customers with high precision. The analyst could then use those insights in campaigns, without waiting weeks for the data science department to deliver a model. In another case, a hospital IT unit applied an AutoML tool to hospital readmissions data and built a risk scoring model for patient readmission, which doctors started using in daily practice to flag high-risk patients – all developed without an outside ML consultancy. These scenarios, increasingly common, show AutoML’s power to put AI development directly into the hands of subject matter experts, effectively democratizing who builds AI solutions.

AutoML in the Machine Learning Landscape

As AutoML adoption grows, it’s important to understand where it fits in the broader machine learning landscape. AutoML is not a silver bullet that eliminates the need for all human involvement, but it has become a core component of modern ML workflows. In practice, AutoML tools are used alongside traditional coding and modeling; data scientists might use AutoML for rapid baselining or to automate parts of the pipeline, while citizen developers rely on AutoML for end-to-end model building. The current landscape can be viewed in terms of the capabilities and maturity of various AutoML solutions, and how they compare to hand-crafted models.

Notably, studies have shown that today's AutoML systems can achieve performance competitive with human experts on many standard tasks. For example, a 2023 empirical benchmark of 12 AutoML tools found that the models generated by AutoML outperformed models that had been manually optimized by researchers for certain classification problems in software engineering. This suggests that for many problems (especially those with well-defined objectives like classification or regression on structured data), AutoML can deliver state-of-the-art accuracy. By systematically exploring a vast space of algorithms and parameters, AutoML sometimes discovers high-performing model configurations that humans might miss. However, the same study also highlighted limitations: none of the evaluated AutoML solutions fully automated every stage of the ML workflow equally well. Some tools excel at model training but do not automate data cleaning or deployment; others automate pipeline steps but offer limited customization. This indicates that the AutoML landscape is still evolving – current tools provide a lot of automation, but often augment rather than replace the human in the loop, especially for tasks like problem formulation, data preparation, and making final trade-offs (accuracy vs. interpretability, etc.).

From a market perspective, the AutoML landscape can be categorized into a few major groups of tools/frameworks:

Open Source AutoML Libraries

These include projects like Auto-sklearn, TPOT, AutoKeras, FLAML, and AutoGluon. They are typically Python libraries that data scientists can install and integrate into notebooks or pipelines. Open source AutoML frameworks often focus on specific tasks (e.g., Auto-sklearn and TPOT target tabular data classification/regression, AutoKeras targets deep learning tasks). Overall, open source AutoML tools are popular with technical users and researchers; they are very capable (often winning or placing highly in public benchmarks) but require the user to orchestrate them within a coding environment.

TIP

In this book, we will use one of these packages, namely AutoGluon to demonstrate the capabilities of AutoML with hands-on examples.

Enterprise AutoML Platforms

Several companies offer commercial platforms that provide an end-to-end AutoML solution with enterprise support. DataRobot and H2O.ai (Driverless AI) are two prominent examples, and cloud vendors like Google (Vertex AI), Microsoft (Azure AutoML), and Amazon (SageMaker Autopilot) also fall in this category. These platforms typically have a GUI, support the entire modeling lifecycle (from data ingestion to deployment), and include features beyond just model training – such as model management, monitoring, and governance. These platforms are also built for deployment maturity: a model can be pushed to production with one click or API call, with automatic scaling. They integrate with databases, cloud storage, and MLOps pipelines, making them attractive for organizations that want a seamless ML pipeline from raw data to deployed service.

When comparing leading AutoML frameworks, a few dimensions stand out as you'll see in the list below:

Accuracy vs. Speed

Tools like AutoGluon and H2O AutoML often prioritize accuracy, sometimes at the cost of longer training or larger ensembles, whereas others like LightAutoML aim to produce decent models very fast (suitable for real-time or limited-resource settings).

Robustness and Resource Use

Some frameworks handle large datasets or noisy data better than others. H2O AutoML has robust performance but can be resource-intensive (needing more memory/compute), which in constrained environments

could be an issue. Others like AutoGluon have shown minimal failure rates even on varied tasks. When scaling up, enterprise platforms usually have built-in checks (and customer support) to handle failures or large jobs.

Feature Engineering and Data Types

AutoML offerings differ in the range of data types and features they handle. Traditional ones like TPOT or auto-sklearn focus on tabular numeric/categorical data. Newer or more advanced ones incorporate text, images, time-series.

Explainability and Compliance

As ML is used in sensitive applications, explainability has become a differentiator (and a necessity, see next subsection on regulations). Many AutoML frameworks now integrate explainable AI (XAI) techniques.

Overall, AutoML has carved out a significant place in the ML landscape: it's no longer an experimental idea but a practical tool used by both novices and experts to accelerate model development. Nearly every major ML platform now includes AutoML capabilities, underlining its importance. It is common to see a workflow where an analyst first uses an AutoML tool to get a quick model and baseline metrics, and then a data scientist might take over to fine-tune or interpret that model. Alternatively, a data scientist might let an AutoML system run overnight to explore hundreds of options, and in the morning they examine the top results and manually refine if needed. In competitions and benchmarks, ensembles produced by AutoML have won or placed well in many challenges (including Kaggle competitions in the past), further proving that automated approaches can stand toe-to-toe with human modelers in many cases. However, experts still play a vital role in guiding AutoML (choosing the right data, defining the problem, and handling nuances that automation can't), and in tackling those custom problems that fall outside of AutoML's current scope.

To summarize the competitive landscape: AutoML frameworks have matured to a point where they can reliably deliver high-quality models across a variety of tasks with minimal human effort. Differences remain in terms of ease-of-use, integration, speed, and special features like explainability and deployment support. Table-based benchmarks (such as OpenML AutoML Benchmark) show multiple frameworks (AutoGluon, H2O, Auto-sklearn, etc.) trading top spots depending on the dataset – indicating that no single tool is universally best, but many are very competent. Organizations often choose based on ecosystem fit: e.g., a company deeply in AWS might lean on SageMaker Autopilot; a team of Python developers might prefer AutoGluon or PyCaret for flexibility; a non-technical group might opt for a managed platform like DataRobot for its GUI and support. The good news is that the AutoML ecosystem is rich and rapidly advancing, driving a healthy competition that continues to improve what these tools can do.

Who Should Use AutoML?

Given the capabilities of AutoML, who benefits most from using it? In practice, AutoML is useful to a wide range of personas – from non-technical domain experts to experienced data scientists – albeit in different ways. A 2024 industry review identified three primary user groups that AutoML empowers:

Software Engineers and Application Developers

These are professionals who need to integrate ML into applications but may not have specialist ML knowledge. AutoML allows them to get a working model without diving deep into algorithm theory.

Citizen Data Scientists (Domain Experts/Analysts)

These are power users in business units – analysts, BI specialists, subject matter experts – who understand the data and problems, but are not formally trained in machine learning. **AutoML is arguably**

designed with this group in mind, to facilitate ML pipeline construction in a low-code or no-code environment.

Data Scientists and ML Engineers

Interestingly, AutoML is also a valuable tool for seasoned ML professionals. Far from replacing them, it can boost their productivity and allow them to tackle more advanced problems. Data scientists use AutoML to automate the tedious parts of experimentation – for example, quickly benchmarking dozens of model types to see what approaches are promising before dedicating time to fine-tune one.

Beyond these groups, organizational leaders and decision-makers indirectly benefit from AutoML as well. AutoML enables faster proof-of-concepts and results, which helps business leaders validate ideas with data. A product manager with an idea for a new ML-driven feature can prototype it with AutoML and demonstrate it within days, helping justify investment. In strategic terms, AutoML should be used by any organization that wants to increase the throughput of ML projects or bring ML capabilities to teams that lack them. This includes small and medium enterprises that might not have a full data science department – AutoML can act as a “force multiplier” for a lone data scientist or an analytics manager, allowing them to deliver value comparable to a larger team.

There are also cases in education and learning: students and learners use AutoML tools to get hands-on with ML concepts without needing extensive coding, which builds intuition and can inspire them to delve deeper. AutoML is thus a great entry point for beginners to experiment with machine learning on real datasets (e.g., try out a dataset in a UI and see results, then inspect which features were important).

On the other hand, it’s worth noting scenarios where AutoML might not be the ideal approach (i.e., who should not use AutoML exclusively). If a problem is highly specialized or the data is very unique (for example, designing a new deep learning architecture for video processing), expert data scientists will hand-craft models beyond the current scope of AutoML.

Similarly, if an organization has ample expert talent and the problem requires maximal optimization or novel research, they might use AutoML only as a reference. In regulated industries, while AutoML provides explainability tools, some institutions might still require a human to validate and adjust modeling choices for compliance or ethical reasons. So, AutoML is not a panacea for all ML tasks – but for the vast majority of routine business prediction and classification problems, and for the broad set of users outlined above, it is an extremely useful technology.

In summary, anyone who needs machine learning insights but lacks either the time or the advanced skills to build models from scratch should consider using AutoML. This spans software developers embedding ML, analysts extracting value from data, and data scientists accelerating their work. By catering to this wide audience, AutoML is helping to cultivate a data-driven culture. Business users can engage directly with ML, and technical experts can deliver more with less effort. The result is that organizations can scale up their ML solutions dramatically – some reports show enterprises building 10X more models per year after deploying AutoML tools – because the power to create models is no longer confined to a small group. As AutoML continues to evolve, we can expect its user base to grow further, including integration with tools that product managers, IT ops, and others use, making AI development truly a collaborative, ubiquitous capability in the enterprise.

AutoML Across Industries: Transforming Business Processes

Automated Machine Learning is rapidly transitioning from a specialized technology to a foundational component of artificial intelligence strategy across a multitude of industry sectors. This widespread adoption is driven by the tangible benefits AutoML offers in terms of efficiency, cost reduction, accelerated insights, and the democratization of advanced analytics. Market forecasts reflect this trend, predicting substantial growth for the AutoML market, with estimates projecting it to reach billions of

dollars within the next few years, indicating significant ongoing investment and adoption. The application of AutoML, however, is highly context-dependent, with different industries prioritizing specific capabilities based on their unique challenges, data landscapes, and regulatory environments.

I'll point out the applications and impact of each in the industries below.

Finance

The financial services industry has been an early and significant adopter of AutoML, driven by the need to process vast amounts of transaction data in real-time and combat increasingly sophisticated fraudulent activities.

- **Applications:** Key use cases include fraud detection, credit scoring and risk assessment, algorithmic trading, and customer analytics.
- **Impact:** AutoML enables financial institutions to build models that can analyze transaction streams in real-time, identifying anomalous patterns indicative of fraud much faster than traditional rule-based systems. A critical challenge in fraud detection is the high rate of false positives generated by older systems, which can be costly and damage customer trust. AutoML-based approaches, such as those implemented using H2O.ai, have demonstrated significant success in reducing false alarms (e.g., by 64%) while simultaneously improving the detection of actual fraud (e.g., by 50-80%). Furthermore, AutoML helps in developing more accurate credit scoring models and optimizing trading strategies. The emphasis on interpretability within some AutoML platforms is also crucial for meeting stringent regulatory compliance requirements in this sector.

Healthcare & Life Sciences

AutoML is making significant inroads in healthcare, promising to enhance diagnostics, personalize treatments, and streamline research.

- **Applications:** Use cases span disease prediction and diagnosis (e.g., **predicting asthma outcomes**, diabetic retinopathy, COVID-19 carriers, **dental conditions**), analysis of medical images (e.g., classifying fundus photographs, segmenting MRI scans), drug discovery, predicting patient readmissions, and **optimizing clinical trial analysis**.
- **Impact:** AutoML tools help researchers and clinicians manage and analyze large, complex, and often multimodal healthcare datasets. They can automate the development of predictive models, potentially improving diagnostic accuracy and enabling earlier interventions. However, healthcare applications face unique challenges, including dealing with highly imbalanced datasets (common in disease prediction), ensuring model generalizability across diverse patient populations, and the critical need for model interpretability and validation before clinical adoption. **Benchmarking studies using medical claims data have shown that while AutoML tools outperform baseline models, significant room for improvement remains, particularly in handling imbalance and achieving high precision-recall performance.** Tools like Google Vertex AI and H2O AutoML are being employed in this domain.

Retail & E-commerce

The retail sector leverages AutoML to optimize operations, personalize customer experiences, and gain a competitive edge in a dynamic market.

- **Applications:** Common applications include demand forecasting, dynamic pricing, inventory management, personalized product recommendations, customer segmentation and churn prediction, and sentiment analysis from customer reviews.
- **Impact:** **AI-driven demand forecasting, often implemented using AutoML frameworks, can significantly improve accuracy compared to manual methods (e.g., from 70% to 90-95%).** This leads to substantial benefits, including reduced overstocking and

carrying costs (by 20-30%), fewer stockouts (reduced by 30-40%), optimized inventory levels, and ultimately, increased sales and customer satisfaction. **Case studies from companies like Bimbo Bakeries (up to 30% forecast error reduction)**, Pacsun, and Belk demonstrate tangible ROI. AutoML also powers recommendation engines and personalization strategies, analyzing customer behavior to deliver targeted offers and improve engagement.

Manufacturing

AutoML is increasingly applied in manufacturing to enhance efficiency, improve quality control, and optimize maintenance strategies.

- **Applications:** Key areas include predictive maintenance (predicting equipment failure based on sensor data), automated quality control (using computer vision for defect detection), supply chain optimization, and process automation.
- **Impact:** **Predictive maintenance models built using AutoML can analyze sensor data (e.g., temperature, rotation speed, torque, tool wear) to anticipate failures**, allowing maintenance to be scheduled proactively, thereby reducing costly unplanned downtime. **In quality control, AutoML powers visual inspection systems that can achieve very high accuracy (e.g., 99%) in identifying product defects, significantly reducing the need for manual inspection (e.g., 80% manpower reduction) and increasing the throughput of quality checks (e.g., 8x increase). Case studies in the automotive industry show AutoML frameworks correlating data from multiple production stages to detect subtle defects early and optimize final vehicle testing procedures, leading to significant time and cost savings.**

Other Sectors

AutoML's reach extends to numerous other industries. Telecommunications companies use it for customer churn prediction and network optimization.

The energy sector applies it for demand forecasting and predictive maintenance of equipment. Agriculture employs AutoML for precision farming to optimize crop yields. Transportation and logistics utilize it for route optimization and supply chain visibility. Government and defense agencies explore its use for cybersecurity threat detection and fraud prevention, while educational institutions apply it for personalized learning paths and administrative efficiency.

The successful deployment of AutoML across these diverse sectors underscores its versatility. However, it also highlights that the choice of tools and techniques must be tailored. Industries like finance and healthcare prioritize interpretability alongside accuracy due to regulatory and ethical considerations, potentially favoring platforms like H2O AutoML or tools with strong XAI integrations. Retail, dealing with massive datasets and prioritizing forecasting, might lean towards scalable platforms or libraries like AutoGluon known for tabular performance. Manufacturing, often involving sensor and image data for real-time decisions, requires robust solutions capable of handling such modalities, potentially leveraging tools like NNI for custom vision models or AutoGluon for multimodal inputs. This context-dependency demonstrates that while AutoML provides powerful automation, strategic implementation requires understanding the specific nuances and priorities of each application domain.

Overcoming Hurdles: Persistent Challenges in AutoML

Despite the rapid evolution and increasing adoption of AutoML, several persistent challenges and limitations hinder its universal application and require ongoing attention from researchers and practitioners. These hurdles often represent the tension between the drive for complete automation and the complexities inherent in real-world machine learning deployment.

Interpretability (The “Black Box” Problem)

Perhaps the most frequently cited challenge is the lack of transparency in models generated by AutoML systems. Many AutoML tools, especially those employing complex ensemble methods or deep neural networks discovered through NAS, produce models whose internal decision-making logic is opaque. This “black box” nature is a significant barrier in high-stakes or regulated domains like healthcare and finance, where understanding *why* a prediction is made is crucial for building trust, enabling debugging, ensuring fairness, detecting biases, and complying with regulations. While the integration of Explainable AI (XAI) techniques like SHAP and LIME is a major focus of third-generation tools providing truly meaningful and reliable explanations for highly complex, automatically generated models remains an active area of research with acknowledged limitations.

Need for Customization vs. Automation

The core value proposition of AutoML is automation, which inherently involves making choices on behalf of the user. However, this can conflict with the needs of expert data scientists or highly specialized applications that require fine-grained control over model architecture, feature engineering steps, or specific hyperparameter settings. Highly automated platforms might lack the flexibility to incorporate specific domain knowledge or implement novel techniques not included in the tool’s predefined search space, potentially leading to suboptimal solutions in niche scenarios.

Data Quality Dependency and Robustness

AutoML systems, like all machine learning approaches, are fundamentally dependent on the quality of the input data. The adage “garbage in, garbage out” holds particularly true; automating the modeling process does not negate the need for clean, representative, and well-prepared data. Poor data quality, including missing values, inconsistencies, errors, or inherent biases, will inevitably lead to poorly performing or unreliable models, regardless of the sophistication of the AutoML tool. Furthermore, ensuring the robustness

of AutoML-generated models—their ability to generalize to new, unseen data and maintain performance in the face of data drift over time—is a critical challenge that requires careful validation and integration with MLOps monitoring practices.

Computational Costs and Resource Intensity

While AutoML aims to save human time, the automated search processes can be extremely computationally demanding. Exploring vast search spaces of algorithms, hyperparameters, and potentially neural architectures requires significant computational resources, often necessitating powerful hardware like GPUs and substantial cloud computing budgets. This cost factor can be prohibitive, particularly for smaller organizations, academic researchers with limited grants, or applications involving extremely large datasets, thereby limiting the accessibility that AutoML aims to promote.

Addressing Bias and Fairness

Ethical considerations are paramount in AI deployment. AutoML systems trained on historical data can inadvertently learn and amplify existing societal biases present in that data. Ensuring fairness and mitigating bias in automated model development pipelines is a complex challenge. While some tools are beginning to incorporate fairness metrics and mitigation techniques (e.g., MLJAR's fairness module), developing robust, automated methods to guarantee fairness across diverse applications remains an open research problem.

Scalability and Efficiency

Although significantly improved over earlier generations, scaling AutoML processes to handle truly massive datasets or extremely complex search procedures (like exhaustive NAS) efficiently can still pose challenges. Balancing the thoroughness of the search with practical constraints on time and resources remains a key consideration.

These challenges highlight that while AutoML offers powerful automation, it is not a panacea. Its successful and responsible adoption requires more than just sophisticated algorithms. It necessitates careful consideration of the application context, robust data governance practices, integration with MLOps for monitoring and maintenance, dedicated efforts towards explainability and fairness, and a realistic understanding of computational costs. Overcoming these hurdles requires a holistic approach, combining technological advancements in AutoML itself with improvements in the surrounding data infrastructure, operational processes, and ethical guidelines.

The Horizon: Future Trends Shaping AutoML

The field of Automated Machine Learning continues to evolve rapidly, driven by ongoing research, technological breakthroughs, and the persistent demand for more powerful, efficient, and trustworthy AI solutions. Several key trends are poised to shape the future trajectory of AutoML:

Synergy with Large Language Models (LLMs) and Foundation Models

One of the most transformative potential developments is the deep integration of AutoML with LLMs and other foundation models. This integration promises a symbiotic relationship:

AutoML for LLMs

AutoML techniques are being explored to automate the complex processes involved in optimizing LLMs themselves, such as prompt engineering, hyperparameter tuning for fine-tuning, and even searching for optimal LLM architectures.

LLMs for AutoML

Conversely, the natural language understanding, reasoning, and generation capabilities of LLMs can enhance AutoML systems. LLMs

could enable more intuitive, conversational interfaces for defining ML tasks, automatically generate code for data preparation and modeling pipelines based on high-level descriptions, assist in automated feature engineering by suggesting relevant transformations based on semantic understanding, and even generate human-readable explanations for model predictions, potentially improving interpretability. Several surveys are beginning to map this emerging landscape.

Foundation Models as Bases

Large pre-trained foundation models (like GPT, BERT, CLIP) can serve as powerful starting points for AutoML. Instead of searching from scratch, AutoML can focus on efficiently fine-tuning these models for specific downstream tasks, leveraging their learned representations to achieve high performance with less data and computation.

Challenges

Despite the excitement, significant challenges remain, including the potential for LLMs to “hallucinate” or generate incorrect code/reasoning, the substantial computational cost of using large models, and inherent ethical concerns regarding bias and reliability.

Next-generation Neural Architecture Search (NAS)

NAS remains a central pillar of AutoML research, particularly for deep learning. Future trends focus on:

Efficiency and Scalability

Developing faster and less computationally expensive search strategies (e.g., gradient-based NAS, one-shot methods where a single “supernet” is trained) to make NAS practical for larger models and datasets.

Broader Applicability

Extending NAS beyond traditional image classification tasks to discover optimal architectures for diverse data types and tasks,

including graph neural networks, transformers for NLP, and time series models.

Improved Search Spaces

Designing search spaces that effectively incorporate prior knowledge and constraints, balancing expressiveness with search efficiency.

Integration with Meta-Learning

Combining NAS with meta-learning to enable faster adaptation of architectures to new tasks.

Maturity of Multimodal Explainable AI (MXAI)

As AutoML systems increasingly handle multimodal data, the need for explainability techniques tailored for these complex inputs becomes critical.

Explaining Cross-Modal Interactions

Developing methods that can elucidate how information from different modalities (e.g., image and text) is combined and influences a model's prediction.

Standardized Evaluation

Establishing robust and widely accepted metrics and benchmarks for evaluating the quality, fidelity, and usefulness of multimodal explanations.

Human-Aligned Explanations

Researching explanation methods that align better with human cognitive processes and provide genuinely insightful, rather than just technically accurate, explanations. The evolution of MXAI is being tracked across different AI eras, from traditional methods to the complexities introduced by generative LLMs.

Continued Democratization via Low-code/No-code

The trend towards more intuitive, graphical interfaces that require minimal or no coding is expected to continue, further broadening the accessibility of AutoML to citizen data scientists, business analysts, and domain experts.

Expansion to Edge Computing and Federated Learning

Practical deployment scenarios are driving the need for AutoML innovations in specialized areas:

AutoML for Edge Computing

Developing techniques to automatically design and optimize lightweight, efficient machine learning models suitable for deployment on resource-constrained edge devices like smartphones, IoT sensors, and embedded systems.

Federated Learning Integration

Incorporating federated learning principles into AutoML frameworks to enable model training across decentralized datasets (e.g., on user devices or different organizational silos) without compromising data privacy.

The future trajectory of AutoML points towards systems that are not only more automated but also more intelligent, interactive, and adaptable. By leveraging breakthroughs in foundation models and LLMs, while simultaneously addressing critical needs for efficiency (NAS, Edge), trustworthiness (MXAI), and privacy (Federated Learning), AutoML is poised to become an even more integral component of the AI landscape.

Summary

In this chapter, we have explored the fundamental concept of Automated Machine Learning (AutoML) and its transformative impact on the field of artificial intelligence. We began by understanding the core drivers behind

AutoML's emergence: the growing demand for machine learning solutions across industries, the persistent data science talent gap, and the need to democratize AI development beyond specialized experts. These factors have created a perfect environment for AutoML to flourish as a technology that automates many of the time-consuming, repetitive tasks in the machine learning workflow.

We've seen how AutoML fits into the broader machine learning landscape, with various open-source libraries and enterprise platforms offering different approaches to automation. We've examined who benefits from using AutoML—from citizen data scientists and domain experts to experienced ML professionals—and how different industries are leveraging AutoML to transform their business processes. Finance, healthcare, retail, and manufacturing are just a few sectors realizing tangible benefits from automated model development.

Despite its advantages, AutoML faces persistent challenges that researchers and practitioners continue to address. Issues around model interpretability, the need for customization versus automation, data quality dependencies, computational costs, and ethical considerations like bias and fairness remain important areas for improvement.

Looking ahead, we identified key trends that will shape AutoML's future evolution. The integration with large language models and foundation models, advancements in neural architecture search, improved multimodal explainability, continued democratization through low-code/no-code interfaces, and expansion to edge computing and federated learning all point toward an increasingly sophisticated and accessible AutoML ecosystem.

Chapter 2. The Rise and Current State of AutoML

A NOTE FOR EARLY RELEASE READERS

With Early Release ebooks, you get books in their earliest form—the author’s raw and unedited content as they write—so you can take advantage of these technologies long before the official release of these titles.

This will be the 2nd chapter of the final book. Please note that the GitHub repo will be made active later on.

If you’d like to be actively involved in reviewing and commenting on this draft, please reach out to the editor at sevans@oreilly.com.

As we saw in Chapter 1, Automated Machine Learning (AutoML) represents a significant paradigm shift within the field of artificial intelligence, moving beyond a collection of tools to embody a comprehensive approach aimed at automating the complex, often laborious, end-to-end process of applying machine learning. The **core objective** of AutoML is to accelerate the development and deployment of machine learning models, making these powerful technologies more efficient and accessible to a wider range of users, including domain experts who may lack deep data science specialization. This automation addresses the inherent complexities and substantial resource requirements—both in terms of time and specialized expertise—that characterize traditional machine learning workflows. An example of such a workflow is shown in Figure 2-1.




Figure 2-1. Machine Learning workflow

The Evolution of Automated Machine Learning

The evolution of AutoML technologies has unfolded through distinct phases, each characterized by significant technological advancements and responses to shifting industry needs. This chapter traces this development through multiple generations, examining the underlying technological drivers, pivotal innovations, and influential platforms that have shaped the field. By understanding this evolution, practitioners can better appreciate the current capabilities and limitations of AutoML systems while anticipating future developments in this rapidly advancing domain. A visual description of this evolution within the larger context of AI is shown in Figure 2-2.

AI took 80+ years to become an overnight sensation




Figure 2-2. A visual history of AI evolution source: <https://blog.planview.com/a-brief-history-of-artificial-intelligence-from-neural-networks-to-chatgpt>

Early Automation (Pre-2010): Laying the Groundwork

The period before 2010 marked the nascent stages of automation in machine learning, characterized by the development of tools and algorithms designed to address specific, isolated challenges within the broader ML workflow. Rather than offering integrated solutions, this era focused on creating algorithmic efficiencies for discrete, often time-consuming, tasks that data scientists routinely encountered.

A primary focus was Feature Selection.

As datasets grew in dimensionality, particularly in fields like bioinformatics and finance, methods to reduce the number of input variables became crucial for model performance and interpretability. Techniques such as forward selection, backward elimination, and LASSO (Least Absolute Shrinkage and Selection Operator) gained traction, helping practitioners identify the most relevant features and mitigate the “curse of dimensionality”. Early efforts in dimensionality reduction also included techniques like **Principal Component Analysis (PCA)**.

Another area of early automation was Hyperparameter Search.

Finding the optimal configuration for a machine learning model’s parameters (which are not learned from data, hence “hyperparameters”) was, and remains, a critical step. Basic methods like **Grid Search**, which exhaustively tries all combinations of specified parameter values, and Random Search, which samples combinations randomly from a defined space, became standard offerings in early ML libraries such as scikit-learn. While straightforward, these methods often proved computationally expensive or inefficient, particularly for complex models with many hyperparameters.

Simultaneously, *Meta-Learning Research* began to emerge within academic circles.

The core idea was to leverage experience gained from solving previous machine learning tasks to accelerate the configuration and training process for new, related tasks. Early studies often utilized benchmark datasets, such as those available from the UCI Machine Learning Repository, to explore how knowledge about algorithm performance on different data types could inform future model selection or hyperparameter settings. This research laid the conceptual groundwork for more sophisticated meta-learning techniques that would become integral to later generations of AutoML systems.

Despite the utility of these individual tools and techniques, their impact was limited by their fragmented nature. Feature selection algorithms, hyperparameter search routines, and meta-learning concepts existed as separate components that required manual integration by skilled data scientists. There was no unified framework or workflow to orchestrate these steps automatically. This phase was fundamentally about developing algorithmic solutions to specific, known pain points within the ML process, driven largely by academic research and the practical need to handle increasingly complex data. It set the stage for the next generation, where the focus would shift towards integrating these disparate components into more cohesive systems.

First Generation (2010-2015): The Dawn of Integrated Systems

The period between 2010 and 2015 witnessed a pivotal shift in the automation of machine learning, moving beyond the automation of discrete tasks towards the development of the first systems that integrated multiple stages of the ML pipeline. This era marked the conceptual birth of modern AutoML, characterized by the unification of algorithm selection and hyperparameter optimization (HPO) into a single, automated process.

The key innovation driving this generation was the formulation and tackling of the **Combined Algorithm Selection and Hyperparameter optimization (CASH)** problem. This approach recognized that the choice of the best machine learning algorithm and the tuning of its hyperparameters are interdependent problems that should be solved jointly for optimal performance. Several pioneering tools emerged, demonstrating the feasibility of this integrated approach:

Auto-WEKA (2013)

Developed by researchers at the University of British Columbia and the University of Freiburg, Auto-WEKA is widely recognized as the first system to explicitly treat the entire CASH problem within a single framework. It operated on the popular WEKA (Waikato Environment for Knowledge Analysis) machine learning library, a Java-based suite of tools.

Auto-WEKA employed **Bayesian optimization** techniques, specifically Sequential Model-based Algorithm Configuration (SMAC), to navigate the complex search space encompassing different algorithms available in WEKA and their respective hyperparameters. This approach provided several key advantages over traditional grid search methods:

- It treated algorithm selection and hyperparameter tuning as a unified optimization problem
- It employed sophisticated modeling of the performance landscape to guide search more efficiently
- It could handle both categorical (algorithm choice) and continuous (hyperparameter values) variables
- It made intelligent trade-offs between exploration (trying new configurations) and exploitation (refining promising ones)

In practice, Auto-WEKA allowed users to specify a dataset and a metric to optimize (e.g., classification accuracy), and the system would automatically explore combinations of preprocessing methods, learning algorithms, and

parameter settings to maximize performance. Empirical evaluations showed it could achieve superior results compared to default configurations or manual tuning by non-experts while requiring significantly less human effort.

While groundbreaking, Auto-WEKA's reliance on the Java-based WEKA library and potential scalability limitations were noted. Its computational efficiency was also challenged by the sequential nature of its optimization approach and the overhead of the WEKA framework.

Hyperopt (2013)

Emerging concurrently, Hyperopt provided a powerful Python library for optimizing complex search spaces, particularly those involving conditional parameters common in machine learning algorithm configuration. It introduced the **Tree of Parzen Estimators (TPE)** algorithm, another effective Bayesian optimization strategy.

Key features of Hyperopt included:

- A flexible framework for defining complex, hierarchical parameter spaces
- Support for conditional parameters (where some parameters only matter if others take certain values)
- The ability to distribute optimization across multiple cores or machines
- Integration with various machine learning frameworks and libraries

Hyperopt gained significant traction within the Python scientific community, especially through its application in the hyperopt-sklearn project, which applied Hyperopt's optimization capabilities to the widely used Scikit-learn library. Comparative studies demonstrated Hyperopt's advantages over simpler Grid Search and Random Search methods in terms of both efficiency and resulting model accuracy.

The impact of Hyperopt extended beyond its immediate application. By providing a general-purpose optimization framework rather than a specialized AutoML solution, it enabled practitioners to apply advanced optimization techniques to a wide range of problems, fostering innovation and experimentation in the field.

SMAC (Sequential Model-based Algorithm Configuration)

While closely associated with Auto-WEKA, SMAC itself is a significant contribution from this era. *As a Bayesian optimization approach often leveraging random forest models to model the performance landscape, SMAC proved highly effective for the high-dimensional, structured, and often discrete search spaces characteristic of the CASH problem. Its success in Auto-WEKA led to its adoption and further development (as SMAC3) in later tools like Auto-sklearn.*

The distinctive features of SMAC included:

- Use of random forests to model the performance landscape, providing robustness to irregularities in the objective function
- Effective handling of categorical parameters common in algorithm selection
- An intensification mechanism to allocate more evaluations to promising configurations
- Support for early termination of poorly performing configurations, improving efficiency

SMAC's importance lay not just in its application to AutoML but in establishing a foundation for efficient optimization approaches in complex, mixed-parameter spaces—a problem fundamental to many aspects of machine learning automation.

First Generation Impact and Legacy

This first generation was heavily influenced by **academic research** and the burgeoning field of **ML competitions**. Platforms like CodaLab hosted challenges associated with major conferences like NeurIPS and ICML, providing crucial proving grounds for benchmarking and refining these early AutoML systems. These challenges not only spurred innovation but also fostered the development of open-source solutions that would become foundational for subsequent progress.

A key takeaway from this period is the establishment of Bayesian optimization (specifically methods like SMAC and TPE) as the dominant technical approach for tackling the CASH problem. This marked a significant advancement over previous brute-force or random search methods. Furthermore, the emergence of both Auto-WEKA (Java/WEKA) and Hyperopt/hyperopt-sklearn (Python/Scikit-learn) highlighted an early divergence between different technological ecosystems, a pattern that continues in various forms today. This generation successfully proved the concept of integrated AutoML, primarily within academic and research contexts, and developed the core optimization machinery that underpins many contemporary systems.

Table 2-1 provides a comparative overview of these influential first-generation tools:

Table 2-1. First Generation AutoML Tools (2010-2015)

Tool Name	Key Technique(s)	Underlying Library/Language	Ease of Use
Auto-WEKA	Bayesian Optimization (SMAC), CASH Formulation	WEKA / Java	★★☆☆☆
Hyperopt	Bayesian Optimization (TPE)	Python (used with Scikit-learn)	★★★☆☆
SMAC	Bayesian Optimization (Random Forest-based)	Language Agnostic	★★☆☆☆

Second Generation (2015-2020): Enterprise Adoption and Cloud Acceleration

The second generation of AutoML, spanning roughly 2015 to 2020, marked a significant transition from academic prototypes to more robust, accessible, and enterprise-oriented solutions. This phase was catalyzed by concurrent advances in computational power, the maturation of cloud computing infrastructure, and a growing industry demand for practical AI applications. AutoML began to move out of the research lab and into the hands of a broader range of practitioners, aiming to democratize capabilities previously limited to experts.

This era saw the emergence and popularization of several influential AutoML frameworks, diversifying the technical approaches and catering to

different user needs:

Auto-sklearn (2015)

Building directly upon the popular Python library Scikit-learn, Auto-sklearn represented a significant step forward within that ecosystem. It extended the CASH problem-solving approach pioneered by Auto-WEKA, employing an updated version of Bayesian optimization (SMAC3). Auto-sklearn brought several substantial innovations to the field:

Meta-learning Capabilities

Auto-sklearn incorporated sophisticated meta-learning, leveraging performance data across numerous previous datasets to “warm-start” the optimization process for new tasks, thereby improving efficiency. When faced with a new dataset, the system would identify similar datasets from its meta-learning database and prioritize algorithms and configurations that had performed well on those similar problems.

Automated Ensemble Construction

The system automatically constructed **ensemble models** from the best-performing individual models found during the search, often leading to improved robustness and accuracy. This ensemble approach combined diverse model types to achieve superior performance compared to any single model.

Comprehensive Preprocessing Support

Auto-sklearn integrated automated preprocessing steps, including imputation of missing values, encoding of categorical features, and normalization of numeric variables, creating more complete end-to-end automation.

Its open-source nature and strong performance made it a popular benchmark and tool, finding applications in European telecommunications and logistics firms. Auto-sklearn consistently demonstrated state-of-the-art

performance on common benchmarks, often matching or exceeding the capabilities of expert-designed pipelines.

However, recent activity on its primary repository suggests potential maintenance concerns. As with many open-source academic projects, long-term sustainability and maintenance emerged as challenges as the field rapidly advanced.

TPOT (Tree-based Pipeline Optimization Tool) (2016)

Developed at the University of Pennsylvania's Epistasis Lab, TPOT introduced a fundamentally different optimization paradigm: **genetic programming**. Instead of optimizing hyperparameters for a fixed model structure, TPOT evolves entire machine learning pipelines, including data preprocessing steps and model choices, represented as tree structures.

Key features of TPOT included:

Evolutionary Pipeline Design

TPOT iteratively combines and mutates pipeline components based on their performance, mimicking natural selection. This approach allowed it to discover complex, multi-step pipelines that might not be obvious to human practitioners.

End-to-End Optimization

The genetic programming approach optimized not just algorithm selection and hyperparameters but the entire workflow, including preprocessing steps and their sequence.

Python Code Export

TPOT gained popularity, particularly in healthcare data analysis, partly due to its ability to export the optimized pipeline as human-readable Python code (using Scikit-learn compatible components), enhancing interpretability and education. This feature allowed users to understand, modify, and integrate the discovered pipelines into larger systems.

Customizable Objectives

The system supported optimization for various objectives beyond accuracy, including balanced accuracy for imbalanced data and custom scoring functions.

While powerful, TPOT's genetic search can be computationally intensive compared to some alternatives, particularly for large datasets or complex search spaces. A successor, TPOT2, is reportedly under development to address some of these efficiency concerns and incorporate newer techniques.

H2O AutoML

Offered by H2O.ai, this platform was explicitly designed with enterprise needs in mind. It provided a scalable and fast AutoML solution supporting both R and Python interfaces. H2O AutoML represented one of the first AutoML frameworks specifically targeting enterprise deployment scenarios rather than academic research.

Key strengths of H2O AutoML included:

Enterprise Focus

The platform prioritized features important in business settings, including scalability, integration with big data platforms, security, and governance capabilities.

Strong Interpretability Tools

H2O AutoML featured robust tools for model explanations and transparency, including model leaderboards, variable importance measures, and partial dependence plots, addressing a critical concern for business adoption.

Data Type Detection

It offered some capability for automatic data type detection, simplifying initial data handling compared to tools like TPOT that required more explicit data preparation.

Stacked Ensembles

H2O AutoML relied on **stacked** ensembles to boost predictive performance, combining predictions from diverse base models through a meta-learning approach.

Technically, H2O AutoML primarily utilized **random search** for hyperparameter exploration rather than the Bayesian optimization approaches common in other frameworks. While theoretically less efficient, this approach offered practical advantages in parallelization and simplicity. Its focus on practical enterprise requirements led to adoption in demanding sectors like fintech and retail.

Google Cloud AutoML (now part of Vertex AI)

Google's entry into the AutoML space leveraged the power and scalability of its cloud platform. It offered distinct, user-friendly tools tailored for specific data types: AutoML Vision, AutoML Natural Language, and AutoML Tables. This domain-specific approach represented a different philosophy from the general-purpose frameworks like Auto-sklearn or TPOT.

A key differentiator was its emphasis on ease of use, providing graphical interfaces that required minimal coding, making it accessible to users without deep ML expertise. Under the hood, Google Cloud AutoML employed sophisticated techniques, including **Neural Architecture Search (NAS)**, particularly for its Vision and NLP offerings. These approaches automated the design of neural network architectures, extending automation beyond traditional ML algorithms to deep learning.

Google's cloud integration provided several advantages:

- **Seamless Scalability:** The platform could handle very large datasets and complex models by leveraging Google's cloud infrastructure.
- **Managed Infrastructure:** Users didn't need to provision or maintain their own computing resources, lowering the barrier to entry.
- **Integration with Google's Ecosystem:** The tools connected with other Google Cloud services for data storage, processing, and deployment.

Its accessibility and domain-specific focus drove widespread adoption, notably in media and marketing analytics, where non-technical teams could leverage sophisticated ML capabilities without requiring specialized expertise.

Second Generation Impact and Legacy

This second generation fundamentally shifted the AutoML landscape in several important ways:

- **Diversification of Optimization Techniques:** The field moved beyond purely Bayesian methods to incorporate genetic programming (TPOT), stacked ensembles (H2O), and neural architecture search (Google), providing users with more options tailored to different problem types and computational constraints.
- **Usability Focus:** A strong emphasis on user experience and accessibility emerged, with tools offering better documentation, simplified interfaces, and integration with popular data science environments.
- **Scalability and Enterprise Features:** Frameworks began addressing production-scale concerns like handling large datasets, meeting enterprise security requirements, and integrating with business systems.
- **Cloud Integration:** Cloud platforms played an indispensable role in this transition, providing the necessary computational resources, managed infrastructure, and accessible interfaces that lowered the barrier to entry for many organizations.

This confluence of factors propelled AutoML from a primarily academic pursuit into a viable tool for broader industry application. Organizations began integrating AutoML into their data science workflows, often using it for rapid prototyping or to address the “long tail” of ML problems that wouldn’t justify dedicated data science resources.

Table 2-2 summarizes the key characteristics of these second-generation platforms.

Table 2-2. Second generation AutoML platforms (2015-2020)

Tool Name	Key Technique(s)	Base Platform/Library	Ease of Use
Auto-sklearn	Bayesian Optimization (SMAC3), Meta-learning, Ensembles	Scikit-learn / Python	★★★★☆☆
TPOT	Genetic Programming	Scikit-learn / Python	★★★★☆☆
H2O AutoML	Random Search, Stacked Ensembles	H2O / Java (R/Python API)	★★★★★☆☆
Google Cloud AutoML	Neural Architecture Search (NAS), Various (per tool)	Google Cloud Platform	★★★★★★

Third Generation (2020-Present): Towards Full-Stack, Multimodal Intelligence

The current, third generation of AutoML technologies, emerging around 2020, represents a significant leap in sophistication and scope. These modern systems extend far beyond the core tasks of hyperparameter

optimization and model selection that defined earlier generations. They aim to automate and integrate larger portions of the machine learning lifecycle, handle increasingly complex and diverse data types, and embed seamlessly within enterprise MLOps (Machine Learning Operations) ecosystems.

AutoML Pipeline Components




Figure 2-3. Typical AutoML workflow

Several leading systems exemplify the capabilities of this generation, as you'll read about here:

AutoGluon (Amazon)

This open-source library, developed by Amazon, has rapidly gained prominence due to its strong out-of-the-box performance and versatility. AutoGluon represents a significant advancement in multiple dimensions:

Multimodal Data Support

Its key differentiator is robust support for **multimodal data**, effectively handling datasets containing combinations of tabular, image, and text features within a single prediction task. This capability addresses a fundamental limitation of earlier AutoML systems that typically focused on a single data type.

Advanced Ensembling Techniques

AutoGluon achieves its impressive performance through sophisticated **ensembling techniques**, particularly a novel multi-layer stacking approach that combines predictions from diverse base models (including deep learning models for text and images alongside classical tabular models).

Deep Learning Focus

Unlike many earlier AutoML tools that concentrated primarily on traditional ML algorithms, AutoGluon places heavy emphasis on deep learning applications, particularly for text and image modalities.

Simplicity and Accessibility

The framework emphasizes ease of use, enabling users to train high-performance models with minimal code—often just a few lines to achieve state-of-the-art results.

AutoGluon has seen adoption in various domains, including e-commerce (recommendations, fraud), insurance underwriting, and customer churn modeling. Tutorials often showcase its multimodal capabilities using datasets like PetFinder, which includes images, text descriptions, and tabular data.

NNI (Neural Network Intelligence – Microsoft)

NNI is another powerful open-source toolkit, positioning itself as a framework for **ML experimentation**. Its focus and architecture differ significantly from AutoGluon, reflecting a different philosophy about AutoML's role:

Neural Architecture Search Focus:

While it supports HPO, NNI's particular strengths lie in facilitating Neural Architecture Search (NAS), model compression, and automated feature engineering. It provides a comprehensive toolkit for discovering optimal neural network architectures.

Experimental Flexibility:

NNI provides flexibility for researchers and practitioners to implement and compare different AutoML algorithms and integrate automated techniques into existing ML infrastructure. Rather than being prescriptive, it enables experimentation with various approaches.

Modular Design:

Its modular design allows users to leverage specific components, like NAS algorithms (e.g., DARTS, ENAS) or feature engineering tools, often via its trial SDK. This component-based approach enables selective automation of specific pipeline elements.

Tuning Framework Integration:

NNI supports multiple tuning frameworks, including TPE, SMAC, Evolution, and custom tuners, allowing users to select the optimization

approach best suited to their problem.

NNI's application in manufacturing for predictive maintenance highlights its utility in complex, real-world scenarios. Its emphasis on neural network optimization makes it particularly valuable for deep learning applications where architecture design is a critical factor in performance.

Google Vertex AI

Representing the evolution of Google Cloud AutoML, Vertex AI is a comprehensive, **unified MLOps platform**. It integrates AutoML capabilities for various data types (Tables, Vision, NLP) alongside tools for custom model training, managed datasets, feature stores, experiment tracking, model registry, prediction endpoints, and pipeline orchestration.

Key advances in Vertex AI include:

End-to-End MLOps Integration

This integrated approach enables true end-to-end ML workflows, from data ingestion to production monitoring. Vertex AI exemplifies the trend toward viewing AutoML not as a standalone tool but as part of a broader MLOps ecosystem.

Multimodal Learning Support

Vertex AI explicitly supports **multimodal learning**, allowing users to build systems that combine insights from different data sources, such as leveraging both text descriptions and image embeddings for product search.

Foundation Model Integration

Recent updates have incorporated access to foundation models (like PaLM, Gemini) that can be fine-tuned or used as feature extractors within AutoML workflows, bridging traditional ML automation with modern foundation model capabilities.

Enterprise-grade Governance

The platform includes robust model governance, lineage tracking, and version control, addressing growing enterprise concerns around responsible AI deployment.

Its broad applicability makes it suitable for diverse sectors, including logistics, energy, and retail. As with earlier Google offerings, Vertex AI leverages the scale and infrastructure of Google Cloud, but now with a more integrated, end-to-end perspective on ML development and deployment.

MLJAR & AWS SageMaker Autopilot

These tools exemplify the trend towards **low-code and no-code AutoML interfaces**. MLJAR, initially a commercial service now available as open-source, offers different operational modes tailored to user goals (e.g., ‘Explain’ for interpretability, ‘Compete’ for maximum accuracy) and uniquely includes a dedicated fairness module to detect and mitigate bias.

AWS SageMaker Autopilot provides an automated solution within the AWS ecosystem, featuring:

- **Code Generation:** Autopilot not only builds models but generates Python code that users can inspect, modify, and learn from.
- **Transparency:** The system provides detailed notebooks explaining its preprocessing, feature engineering, and model selection decisions.
- **Integration:** Tight coupling with the broader SageMaker ecosystem for deployment, monitoring, and lifecycle management.

Such tools further lower the barrier to entry, making AutoML accessible to business analysts, educators, small businesses, and agile enterprise teams who may lack extensive coding skills.

Third Generation Key Capabilities

Beyond specific tools, the third generation is defined by several advanced capabilities becoming increasingly mainstream:

Multimodal Learning

The ability to natively handle and fuse information from diverse data types (tabular, text, image, potentially audio/video) is a hallmark of modern systems like AutoGluon and Vertex AI. This is crucial as real-world problems rarely involve only one type of data.

Technical approaches include:

- **Specialized Encoders:** Using optimized architectures (e.g., Transformers for text, CNNs/Vision Transformers for images) for different data types
- **Fusion Mechanisms:** Sophisticated techniques for combining information from different modalities
- **Representation Learning:** Automatic learning of joint representations across data types
- **Transfer Learning:** Leveraging pre-trained models for each modality and fine-tuning them for specific tasks

This multimodal capability represents a step-change in AutoML's applicability to complex, real-world problems where information exists across different formats and sources.

Explainability Integration (XAI)

Recognizing the critical need for transparency, third-generation tools increasingly embed XAI techniques directly into their workflows. Methods like SHAP (Shapley Additive exPlanations) and LIME (Local Interpretable Model-agnostic Explanations) are commonly used to provide both global feature importance scores and local explanations for individual predictions.

This integration serves multiple purposes:

- **Building User Trust:** Helping users understand and trust model decisions
- **Debugging and Improvement:** Enabling more effective model debugging and refinement
- **Regulatory Compliance:** Meeting growing regulatory requirements for model transparency
- **Ethical AI:** Supporting ethical assessment and bias detection in model behavior

Platforms often generate automated interpretability reports or dashboards, making explanation a standard part of the AutoML workflow rather than an optional add-on.

MLOps Integration

Modern AutoML is no longer a standalone process but an integral part of the broader MLOps lifecycle. AutoML pipelines are designed to integrate with CI/CD (Continuous Integration/Continuous Deployment) practices, version control systems (GitOps), feature stores for managing reusable features, and monitoring tools for tracking model performance and drift in production.

This integration addresses several key challenges:

- **Deployment Gap:** Reducing friction between model development and production deployment
- **Governance:** Supporting model governance, versioning, and compliance requirements
- **Monitoring:** Enabling ongoing performance tracking and drift detection

- **Retraining:** Facilitating automated retraining when performance degrades

Platforms like Vertex AI epitomize this trend by offering a unified environment for managing the entire journey from data to deployed, monitored model.

Neural Architecture Search (NAS)

NAS continues to be a core research area and a key component of advanced AutoML systems, particularly those focused on deep learning (e.g., NNI, cloud platforms). The goal remains the automatic discovery of optimal neural network structures for specific tasks.

As of today, research direction seems to focus on:

- **Search Efficiency:** Improving the efficiency and effectiveness of the search process (balancing exploration vs. exploitation)
- **Performance Estimation:** Developing better strategies for estimating architecture performance without full training (reducing the need for costly full training runs)
- **Hardware-aware NAS:** Optimizing architectures not just for accuracy but for deployment constraints like latency and memory usage
- **Specialized Architectures:** Discovering architectures optimized for specific domains and data types

The maturation of AutoML into end-to-end platforms capable of handling multimodal data and integrating seamlessly with MLOps signifies a major shift. The focus is now on providing practical, scalable, and governable AI solutions for complex, real-world enterprise problems, moving well beyond the initial goals of simply automating model selection and tuning. The incorporation of XAI and the availability of low-code interfaces further enhance the practicality and accessibility of these powerful technologies.

Table 2-3 contrasts key third-generation offerings.

Table 2-3. Third Generation AutoML Tools (2020-Present)

Tool/Platform Name	Key Technique(s)	Primary Approach	Ease of Use	Performance
AutoGluon	Multimodal Learning, Advanced Ensembling	Open Source Library	★★★★☆	★★★
NNI	NAS, HPO, Model Compression	Open Source Toolkit	★★☆☆☆	★★★
Vertex AI	Unified MLOps, Integrated AutoML	Cloud Platform (GCP)	★★★★★	★★★
MLJAR	Low-code Interface, Fairness Module	Open Source Platform	★★★★★	★★★
SageMaker Autopilot	Code Generation, Transparency	AWS Cloud Service	★★★★☆	★★★

Summary

The evolution of AutoML across its three generations illustrates a remarkable journey from specialized academic tools to comprehensive enterprise platforms. This progression has been characterized by expanding scope, increasing accessibility, and deeper integration with the broader machine learning ecosystem.

The first generation (2010-2015) established the foundational concept of automating algorithm selection and hyperparameter tuning as a unified optimization problem. Tools like Auto-WEKA, Hyperopt, and SMAC demonstrated the feasibility of this approach and introduced powerful Bayesian optimization techniques that remain influential today.

The second generation (2015-2020) brought AutoML into the mainstream, with frameworks like Auto-sklearn, TPOT, H2O AutoML, and Google Cloud AutoML offering more robust, user-friendly solutions that began finding traction in enterprise environments. This period saw significant diversification in technical approaches, from meta-learning and genetic programming to neural architecture search and ensemble methods.

The current third generation (2020-present) has expanded AutoML's scope to address end-to-end machine learning workflows, multimodal data, and seamless MLOps integration. Systems like AutoGluon, NNI, and Vertex AI exemplify this broader perspective, offering sophisticated capabilities for handling diverse data types and supporting the complete lifecycle from development to deployment and monitoring.

Looking ahead, several trends suggest the future direction of AutoML development:

- **Foundation Model Integration:** Increasingly tight coupling between AutoML and foundation models (like large language models and multimodal models)
- **Sophisticated XAI:** More advanced explainability techniques to address the growing demand for interpretable AI
- **Further Democratization:** Continued emphasis on accessibility through natural language interfaces and low-code/no-code

platforms

- **Specialized Vertical Solutions:** Domain-specific AutoML solutions optimized for particular industries and use cases

The AutoML landscape continues to evolve rapidly, with each new generation building upon and expanding the capabilities of its predecessors. As these technologies mature, they promise to make sophisticated machine learning capabilities more accessible, efficient, and practical for a growing range of users and applications. The journey from early academic experiments to enterprise-ready platforms reflects not just technical progress but a broader democratization of AI that has significant implications for how organizations leverage data and analytics in the years ahead.

Looking Ahead: The Critical Role of Hyperparameter Optimization

As we've traced the evolution of AutoML from its early beginnings to today's sophisticated multimodal, MLOps-integrated platforms, we've repeatedly encountered hyperparameter optimization (HPO) as a foundational component (Figure 2-4). This critical capability was present in the earliest AutoML systems and remains central to the most advanced contemporary tools.




Figure 2-4. Current set of optimization approaches

In the next chapter, we'll take a deep dive into hyperparameter optimization, addressing several fundamental questions:

- **How do different HPO algorithms balance exploration and exploitation?** We'll examine the mathematical foundations of Bayesian optimization, evolutionary algorithms, and multi-fidelity approaches, comparing their efficiency and effectiveness across different problem types.
- **What makes hyperparameter optimization particularly challenging in deep learning contexts?** We'll explore the unique challenges of optimizing neural network hyperparameters, including the high dimensionality of the search space, complex interactions between parameters, and the computational expense of evaluations.
- **How can we incorporate domain knowledge into the HPO process?** We'll investigate techniques for incorporating prior knowledge, constraints, and domain-specific insights to make the optimization process more efficient.
- **What are the practical trade-offs between different HPO approaches?** We'll provide pragmatic guidance on selecting and configuring HPO methods based on computational resources, problem characteristics, and accuracy requirements.
- **How is HPO evolving in the context of foundation models?** We'll examine emerging approaches for efficiently fine-tuning large pre-trained models, where traditional HPO methods may be prohibitively expensive.

By understanding the nuances of hyperparameter optimization—perhaps the most mature and fundamental component of AutoML—practitioners will gain insights applicable to the entire AutoML workflow. The next chapter will equip you with both theoretical understanding and practical techniques

to implement effective hyperparameter optimization strategies for your specific ML challenges.

About the Author

Dr. Kerem Tomak is a seasoned AI and data science leader with extensive experience implementing machine learning solutions at scale. As Global Chief Data & Analytics Officer at Decathlon and former executive at companies like ING, Commerzbank AG, and Google, he has led the design and deployment of automated, AI-powered systems across finance, retail, and tech. With a Ph.D. in Management Information Systems from Purdue University, multiple patents in machine learning, and a history of driving enterprise AI adoption, Dr. Tomak brings deep expertise to the evolving field of AutoML.